

# Introductory concurrency

Mark Armstrong

November 27, 2020

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Some brief background</b>	<b>2</b>
2.1	Concurrency vs. parallelism . . . . .	2
2.2	Process vs. thread . . . . .	2
2.3	Race condition vs. deadlock vs. starvation . . . . .	3
<b>3</b>	<b>Creating threads with futures in Clojure</b>	<b>3</b>
3.1	Representing trees in Clojure . . . . .	4
3.2	Finding the maximum of a tree . . . . .	4
3.3	Using futures . . . . .	5
<b>4</b>	<b>Forking processes in Ruby</b>	<b>5</b>
<b>5</b>	<b>Justifying our example problem</b>	<b>5</b>

## 1 Introduction

Here we briefly explore (some of) the techniques to take advantage of concurrency and parallelism in the languages we have used throughout the course.

Our context for this discussion will be the problem of computing the maximum element of an (unsorted) tree of integers. This naturally leads to a divide and conquer approach, which can easily take advantage of concurrency, while also avoiding the issues that can arise from concurrent execution. We do give a brief discussion of these issues before starting our exploration.

We justify the use of this problem to frame our exploration with a discussion at the end of the notes.

## 2 Some brief background

The following should be review from other courses, most notably from “Concurrent Systems”. If you are not familiar with these topics, this will be a sufficient introduction for our purposes here.

### 2.1 Concurrency vs. parallelism

To define concurrency and parallelism, we must refer to “tasks being executed”.

Two tasks A and B are said to be *concurrently executed* if the execution may be carried out in any of these orders:

- Part of A is executed, then the remainder of A and B are executed concurrently.
- Part of B is executed, then the remainder of B and A are executed concurrently.
- Part of A and part of B are executed *at the same time*, then the remainder of A and B are executed concurrently.

Note that tasks can be run concurrently *even if* it is not possible for the underlying system to execute two tasks at the same time.

On the other hand, two tasks are said to be *executed in parallel* if their execution takes place (at least partly) at the same time.

Parallelism then requires hardware support to run two tasks at the same time. In the past, this implied distributed systems, but on modern computers, the presence of multiple *cores* of the processor allow for local parallel computation.

As the notion of concurrency is more general, we usually talk about concurrent execution, unless we specifically mean that tasks are executed at the exact same time.

### 2.2 Process vs. thread

A *process* is an instance of a program being executed on a machine. The machine, usually through an operating system, grants each process its own memory space, and handles the scheduling of the execution of this and all other processes.

A *thread* is a single sequence of instructions to be executed by the machine. Threads are not granted their own memory space by the machine (but the process may handle assigning memory to threads.)

A single process may be made up of several threads; so long as the process is being executed, there is at least one thread (of execution.)

### 2.3 Race condition vs. deadlock vs. starvation

Race conditions, deadlock and starvation are all classes of error that can occur when using concurrency.

A *race condition* refers to an error that arises when ...

A *deadlock* refers to an error that arises when ...

A task is said to *starve* when ...

## 3 Creating threads with futures in Clojure

Clojure has support for several approaches to concurrency, including but not limited to:

- a STM (Software Transactional Memory) system,
- an *agent* system (somewhat related to the *actor model*), and
- an *atom* system which avoids race conditions regarding changes to data.

Refer to the [documentation](#) for more information about these.

For today, we are well served with a simpler approach: the `future` macro. A sequence of expressions wrapped in a `future` are evaluated in a new thread.


```
(future (Thread/sleep 4000) ;; 4 seconds
  (println "Computation on this thread paused and this
    ↪ printed after 4 seconds."))

(println "Computation on the main thread continued and this
  ↪ printed immediately.')
```

Before we look closer at futures, we must discuss how we represent these trees in Clojure. Then we will show a single-threaded implementation of the maximum function, and finally rework it into an approach using threads via futures.

### 3.1 Representing trees in Clojure

Quite simply, we will consider *nested sequences* to be trees. We will use the convention (at least for these notes) that the first element of the list is the *label* of a node, and the next two elements of the list gives the *children* of that node.



For instance, the tree  
can be written

```
'(1 (2) (3 (4) (5)))
```

Note that in the case that a tree has no children, we omit the rest of the list.

### 3.2 Finding the maximum of a tree

```
(defn maximum-tree
  "Find the maximum element of a tree of integers `t`.
  `t` is assumed to be a list whose first element is an integer
  (the label of the root) and whose remaining elements are lists
```

```

representing trees containing integers."
[t]
(if
  (empty? t) ##-Inf
  ;; Use list deconstruction to separate the `label` and the
  ↪ list of `children`.
  (let [[label left right] t]
    (max label (maximum-tree left) (maximum-tree right)))))

(maximum-tree [1 [2] [3 [4] [5]]])

```

### 3.3 Using futures

:TODO:

## 4 Forking processes in Ruby

:TODO:

## 5 Justifying our example problem

:TODO: